




# Toward Reliable Biodiversity Dataset References


Michael Elliott<sup>1</sup>, Jorrit H. Poelen<sup>2</sup>, José A.B. Fortes<sup>1</sup>



**1 Advanced Computing and Information Systems Laboratory  
(ACIS), University of Florida, Gainesville, Florida, USA**  
**2 400 Perkins St Apt 104, Oakland, California, USA**

# Toward Reliable Biodiversity Dataset References

## Abstract

No systematic approach  has yet  adopted  to reliably reference and provide access to our digital biodiversity datasets. Based on accumulated evidence, we argue that location-based identifiers such as URLs are not sufficient to ensure long-term data access. We introduce a method that uses dedicated data observatories to evaluate long-term URL reliability.

After taking periodic inventories from March through October 2019 of the data served by major biodiversity aggregators, including GBIF, iDigBio, DataONE, and BHL, we found that, for each network, 5% to 43% of registered  URLs were intermittently or consistently unresponsive, 0% to 63% produced unstable content, and 13% to 76% became either unresponsive or unstable over the period of observation.

We propose the use of  cryptographic hashing to generate content-based identifiers that can reliably reference datasets. We show that content-based identifiers facilitate decentralized archival and reliable distribution of biodiversity  datasets. Among other potential benefits, our proposed identifiers enable reliable dataset citations and long-term reproducibility of biodiversity studies.

**Keywords**— Biodiversity, Ecological Informatics, Information Systems, Information Retrieval

# Introduction

23

Over the course of hundreds of years, naturalists and biologists have  
systematically collected physical evidence from an ever-changing natural  
world. Through well-established protocols and institutional support, many  
of these natural history collections have withstood the ravages of time  
(Davis and Schmidt 1996, Hortal et al. 2015). Records that describe these  
carefully collected specimens are now made available digitally through  
online search indices, registries, and data archives (Page et al. 2015). The  
increased availability of digital natural history records helps work toward  
Charles Elton’s realization that ecosystems can only be fully understood  
when we “provide conceptions which can link up into some complete  
scheme the colossal store of facts about natural history which has  
accumulated up to date in this rather haphazard manner” (Elton 1927).  
So far, various initiatives have succeeded to provide comprehensive  
aggregate views from previously scattered natural history record siloes  
((alias?), GBIF 2019, Matsunaga et al. 2013, Michener et al. 2011,  
Rinaldo and Norton 2009). However, we show that these aggregate views  
are subject to change as their underlying digital source data changes or  
becomes inaccessible. Although efforts have been made to keep track of  
changes in dataset networks, such as using version numbers and last  
modified dates (Robertson et al. 2014, Wieczorek et al. 2012) and periodic  
archival (Costello et al. 2013), we are not aware of the adoption of any  
systematic approach to preserve the accessibility as well as longevity of our  
digital natural history record and derived datasets. We have collected  
evidence that, despite hundreds of years of experience in preserving our  
physical natural history records, we are currently faced with a growing

body of digital data that changes daily and can disappear with the push of  
a button.

Our scholarly record is stitched together by an intricate web of  
associations between scientific studies and the datasets on which they are  
based. These associations are made explicit through citations that can be  
used to reconstruct a study's context and provide the chain of evidence  
that support its claims (Garfield et al. 1964). In the pre-internet era, the  
lookup of cited references required access to one or more of the many  
academic libraries in the world. With the rise of internet accessible  
scientific publications, authors and readers access these references by using  
a networked device to download content from publication websites. This  
means that researchers are increasingly citing online works to support their  
claims. Because the citation format of online works typically documents  
only when (e.g., 2019-10-01) and where (e.g., <https://doi.org/10.123/456>)  
the referenced work was accessed by the author (DataONE 2012, GBIF  
2019, iDigBio 2016), the reader expects the web-accessed resource to  
remain accessible and unaltered via this single web location. Readers may  
attempt to find a version of the works referenced by searching online data  
networks for the matching author and title, but there is no guarantee that  
information found this way will be exactly the same as what was originally  
referenced. Any reference that does not allow readers to find the referenced  
work fails to satisfy the FAIR principle of findability: "F1. (meta)data are  
assigned a globally unique and eternally persistent identifier" (Wilkinson  
et al. 2016). Our study is not alone in providing evidence that suggests  
that networked, location-based access to digital objects is an unreliable  
mechanism for providing continued access to the unaltered original work

(Klein et al. 2014, Vision 2010). Unless we change the way we preserve and  
cite our digital scholarly works, our physical records stored in libraries and  
museums around the world are likely to outlast our digital ones.

## Problem Characterization

We show that the current practice of using Uniform Resource Locators  
(URLs) (Berners-Lee et al. 1994) to reference online biodiversity datasets  
provides no guarantee of long-term data accessibility. Readers who  
encounter references that use URLs as dataset identifiers cannot be certain  
that the referenced dataset will continue to be accessible and in its original  
form. This uncertainty might be cause for alarm for researchers because,  
over time, the integrity of the scholarly record itself is damaged when  
existing references become unreliable due to the loss of access to the data  
they reference. When data access is lost, documented research results may  
become impossible to reproduce and the justification for conclusions or  
hypotheses that relied on lost results may be undermined.

The current practice of relying on URLs to locate and identify  
referenced data is hazardous due to their demonstrated risk of link rot and  
content drift (Klein et al. 2014). Link rot occurs when a URL, or link, that  
previously responded to queries can no longer be reached. This can  
happen, for example, due to temporary outages, URL retirement, or URL  
migration. A link exhibits content drift when a query to the link provides  
content that is different from the content it provided in the past. The  
extent of content drift can vary; content may have received only minor  
edits with no changes in semantics, or it may reference a different entity  
altogether. When a single URL is used to locate data that may change

over time, access to any particular version of the data is likely to be 100  
 should be removed. We show that, in the event of link rot or content drift, any 101  
 existing references that point to the affected URL may become unreliable. 102  
 In one study on the *Geophysics* journal, it was reported that 40% of links 103  
 (URLs) to supplemental materials became unavailable due to link rot 104  
 within one year of publication (Vision 2010). Another study (Klein et al. 105  
 2014) confirmed that as many as one in five articles in the *Journal of Science, 106  
 Technology, and Medicine* provide references that exhibit either link rot 107  
 and content drift and refer to the existence of either as “reference rot”. As 108  
 discussed in the results section of this paper, we have found that the use of 109  
 URLs to reference biodiversity datasets according to existing biodiversity 110  
 network citation guidelines (DataONE 2012, GBIF 2019, iDigBio 2016) 111  
 can lead to unreliable dataset references as a result of reference rot. The 112  
 information systems used by major biodiversity data networks, such as 113  
 DataONE, GBIF, and iDigBio, rely on data curators and institutional 114  
 repositories to maintain active dataset URLs, and aggregate the data 115  
 found at those URLs for distribution in response to user queries. If a data 116  
 curator modifies, relocates, or stops serving a particular dataset, it may 117  
 become impossible to retrieve the original dataset and the integrity of the 118  
 data network will suffer as a result. 119

In this paper, we propose a methodology for measuring the existence of 120  
 link rot and content drift in online data networks, then provide 121  
 experimental results that confirm the existence of both link rot and 122  
 content drift across all of the biodiversity data networks we considered, 123  
 including BHL, DataONE, iDigBio, and GBIF. Finally, we propose a 124  
 method for referencing and serving biodiversity data in a way that works 125


toward satisfying the Findable, Accessible, Interoperable, and Reusable (FAIR) principles (Wilkinson et al. 2016).

## Methodology

Although it has been demonstrated that reference rot does occur when URLs are used for referencing scientific works (Klein et al. 2014, Vision 2010), we are not aware of any prior studies that provide quantitative evidence that reference rot occurs specifically in biodiversity data networks.

We set out to quantify the extent of reference rot in biodiversity data networks. Because reference rot occurs in the scope of individual data references, and references to digital datasets rely on URLs to locate the data, we begin by introducing terminology for characterizing the reliability of a URL according to how often it exhibits link rot and content drift.

## URL Reliability

We assume that the URLs used to reference biodiversity datasets are expected to resolve to an Internet Protocol (IP) (Postel 1981) address via the Domain Name System (Mockapetris 1987). If a web server is accessible at the resolved IP address, a query (i.e., HTTP get request) to that address over the Hypertext Transfer Protocol (HTTP) will return a response code and, in some cases, associated content (Berners-Lee et al. 2005). We classify the reliability of a URL according to the content, or lack of  that it provides over successive queries. If a query to a URL is unsuccessful, we say that link rot has occurred. However, if a successful response is received but the retrieved content is different from the content

retrieved by previous query, we say that content drift has occurred. 149

Monitoring URLs in this way allows us not only to determine whether link 150

rot and content drift occur, but also to capture their long-term behaviors. 151

For example, one URL that has exhibited link rot might have failed to 152

respond only once, whereas another might have become **repeatedly** 153

unresponsive. **Likewise, one URL might exhibit content drift less frequently** 154

**than another whose contents change rapidly.** Furthermore, various 155

combinations of link rot and content drift behavior may indicate that one 156

URL is more reliable than another, even though both exhibit reference rot. 157

We label URLs with sets of reliability indicators according to their link 158

rot and content drift behaviors. The defined reliability indicators are 159

differentiated by the degree of link rot and **content drift** observed over a 160

series of queries to the URL at different points in time. We characterize 161

the responsiveness of a URL according to how often it exhibits link rot: 162

- Unresponsive: the link has failed to respond to one or more queries 163
- Responsive: the link has responded to all recorded queries 164

We characterize the stability of a URL according to how often it 165

produces different content from one query to the next: 166

- Unstable: the content that the link points to sometimes changes 167
- Stable: the content that the link points to never changes 168

We characterize the overall reliability of a URL according to both its 169  
responsiveness and stability: 170

- Unreliable: the link does not always provide the expected content; it 171  
is either unresponsive, unstable, or both 172



- Reliable: the link always provides the expected content; it is both responsive and stable

Before we can determine the reliability of any given URL, we must first monitor its behavior over time by documenting how it responds to periodic queries. For the context of biodiversity, we consider the case in the content that a URL produces is a dataset.

## The Data Collection Process

We suggest that digital dataset collection practices have some analogies to well-established physical specimen collection procedures (see figure 1) (Poelen 2019g). If datasets are considered analogous to specimens, then the URLs that locate datasets are analogous to the physical locations of specimens in the natural world; they are where digital datasets were originally found, but not where they should be preserved. Once found, physical specimens are collected by hand; similarly, digital datasets are downloaded by querying their URLs. Once a specimen is collected and deposited to a safe, accessible repository, a record is kept that documents what the specimen is in addition to when, where, and by whom it was collected.

The same can be done for downloaded datasets. When a dataset is downloaded, a record can be kept that details the URL that was queried, the time of query, and who (e.g., a human or software agent) issued the query that initiated the download event; we refer to this record as the dataset's provenance record (Pasquier et al. 2017). Additionally, the dataset itself should be stored in a safe, accessible dataset archive. The final step in the collection process is to link the actual preserved specimen

to its corresponding record (see figure 1(a)) via an assigned unique  
identifier.

The identifiers assigned to datasets (or provenance data) must differ  
only if their contents differ. This can be achieved by deriving the  
identifiers from the content of their respective datasets. Furthermore, the  
identifier must be unique to the dataset; a dataset will always be assigned  
the same identifier and no two datasets (including different versions of a  
dataset) can share an identifier. With these restrictions in place, it is  
possible to reliably link each recorded dataset to its provenance record  
without the need for an intermediate index. That is, the derivation of the  
content-based identifier from a given dataset can be performed by anyone,  
anywhere, and at any time, without the need to consult some central  
authority (Paskin 1999). Cryptographic hashing is one such method for  
producing content-based identifiers which are both content-derived and  
unique. A variety of cryptographic hashing algorithms exist which receive  
some digital file as input and uniquely encode its contents into a  
fixed-length series of bits called a “hash”. We use hashes generated by the  
SHA-256 algorithm (NIST 2001) as unique content-based identifiers.

## Data Collection Over Time

By establishing a dedicated data observatory that implements the  
collection process we have described, we can build a history for each  
observed URL to capture its long-term reliability. Such an observatory  
periodically queries the URLs listed in data network’s URL registry and  
produces for each URL two complementary parts: 1) an archived copy of  
the response to the corresponding query, whether it was a dataset, an error

code, or no reply at all, and 2) a record of its provenance, including the URL itself, the current date, and a content-based identifier of any dataset received. Successive provenance records can be aggregated to construct comprehensive histories for both datasets (when and where they were found) and URLs (which datasets they located over a series of queries over time).

The constructed URL histories can be analyzed to determine whether a link was ever broken, when it was broken, and whether it became responsive again. The logs also identify the content (or lack of) that a URL located each time it was queried. Any change in the content identifier from query to the next indicates a change in the content of the dataset. These link breakages and content changes correlate to link rot and content drift, respectively, and allow us to determine the responsiveness, stability, and reliability of each URL over time.

## Data Network Reliability

Our method for monitoring the behavior of a single URL over time can be applied to observe all URLs registered in a biodiversity data network. We also extend the idea of URL reliability to entire data networks and propose that the overall reliability of a data network can be evaluated by monitoring the long-term reliability of each individual URL in the network *ex post*. Whereas we rigidly label individual URLs with binary indicators of responsiveness, stability, and reliability, we grade data networks according to the percentage of registered URLs that are assigned each of the reliability indicators. For example, if a data network contains three distinct URLs and we find that only two out of the three are reliable, then

we say the data network is 67% reliable.

248

## Experiment

249

The Preston biodiversity dataset tracker (Jorrit Poelen and Patel 2018)

250

implements mechanisms for monitoring data networks as we have

251

described. It allows users to deploy a data network observatory

252

systematically observes the entire set of URLs registered in the network,

253

queries each URL for data, then documents data collection and archives

254

the results. All crawl activities, the queries they issue, and the results they

255

produce are meticulously recorded in a string of provenance logs.

256

We deployed several Preston observatories which periodically queried

257

the registered dataset URLs listed by Biodiversity Heritage Library (BHL),

258

Data Observation Network for Earth (DataONE), Global Biodiversity

259

Information Facility (GBIF), and Integrated Digitized Bio Collections

260

(iDigBio). Each of these networks provides online registries of URLs that

261

locate the data in the network. The registered URLs for DataONE, GBIF,

262

and iDigBio were queried monthly from March 2019 through October 2019.

263

BHL was queried monthly from May 2019 through October 2019. The logs

264

taken by each of these observatories describe the URL queries and their

265

results, which were processed to produce the results that follow. A

266

observatory was constructed by aggregating the queries of the data

267

network observatories.

268

## Results

Breakdowns of the overall reliabilities of the data networks are provided in table 1. Results are listed as percentages and total counts of URLs in the data network that were assigned each reliability indicator. When analyzing the recorded results of queries to URLs in each data network over a period of several months, we found that, for each individual network, 5% to 43% of registered URLs were intermittently or consistently unresponsive, 0% to 63% produced unstable content, and 13% to 76% became either unresponsive or unstable over the period of observation.

We found that 30% of URLs observed across the networks became unreliable at some point over the period of March 2019 through October 2019. Of those unreliable URLs, 48% were unstable, 22% became consistently unresponsive, and 70% were at best only intermittently responsive. For 5% of successful queries, the URL failed to respond to the next query. For 4% of successful queries, the URL provided different content the next time it responded when queried.

The changes in reliability over time for each network are visualized in figure 2. Note that because we have defined reliable URLs to be those considered both responsive and stable, they always represent the smallest fraction of URLs in table 1, figure 2, and figure 3. Figure 3 visualizes the cumulative growth of biodiversity data networks during their periods of observation. This growth is illustrated with two metrics: the total number of unique URLs ever registered in each network and the total number of unique contents that have been downloaded from the network at each sampled point in time.

The behaviors of the distributions over time of responsive, stable, and

reliable URLs vary notably between data networks. Reasons for these  
 differences might be inferred when cross-examining table 1 and figures 2  
 and 3. For example, although BHL scored relatively low in responsiveness  
 due to frequent link rot, the content that it does provide is more stable  
 than all other networks because content drift within BHL is relatively rare.  
 Conversely, although iDigBio is relatively responsive, it has low stability  
 because the network's near-constant content growth far outpaces its URL  
 growth. GBIF's behavior was characterized by large sporadic swings; a  
 mass URL migration of over 14,000 Plazi-hosted datasets occurred in May,  
 introducing thousands of new URLs over a short period of time, while over  
 31,000 URLs (60% of URLs that responded to queries that month)  
 suddenly changed contents in October. Even the most reliable network,  
 DataONE, shows a clear downward trend in all three categories, with 13%  
 of URLs becoming unreliable over a period of just seven months.  
 Additionally, DataONE's growth curves indicate that there are far fewer  
 unique contents than unique URLs; this suggests two possibilities: either  
 much of DataONE's URL population is unresponsive, or DataONE lists  
 multiple URLs for many of its datasets. Because DataONE has been  
 shown to be highly responsive, it could be the case that many distinct  
 URLs refer to the same datasets. It's also worth noting that the June and  
 September spikes in BHL's unresponsiveness were largely due to URLs  
 that failed to respond in those particular months but eventually did respond  
 to future queries.

## Sources of Potential Numerical Error

318

We expect that the URL reliability counts generated for the figures and tables are lower than their actual values. When we qualified URLs as being reliable, responsive, and stable, we could not be certain that links did not briefly become unresponsive or change content during the month-long periods between queries. It is therefore likely that some cases of link rot and content drift were not reflected in the results. Additionally, we only query URLs that the data networks list in their dataset registries; this means that, after URLs are removed from a network’s registry, we did not detect subsequent instances of reference rot. Therefore, our results represent a very optimistic upper bound on URL and network reliabilities.

The results for DataONE and GBIF in figure 2 are sometimes skewed due to the pagination method that the networks use to supply users with their dataset registries. Registry pages contained set amounts (e.g., 20) of URLs and represent small slices of the actual data network registry. For registries that use pagination, the observatory would keep querying for registry pages until reaching the page or failing to reach the end. For instance, GBIF’s URL and dataset totals in March 2019 (see figure 3(c)) are low because an early query to a GBIF registry page was not answered and, consequently, the URLs of registry pages that should have followed were not discovered. Similar events happened for both the GBIF and DataONE observatories at later points in time, potentially overestimating the reliability of the data network.

In an effort to minimize artificial link rot due to internet access issues in our local network, we deployed the Preston observatories in a large commercial data center in Germany.

## Discussion

We have shown that the reliability of URLs decreases over time in all of the major biodiversity data networks that we monitored. If current trends continue, the extent of reference rot will worsen. Systematic changes in the way we preserve and reference data are needed to reverse these trends and improve the longevity and long-term integrity of the biodiversity data record. Before we propose such changes, it's necessary to first understand why URLs are proving to be ill-suited for referencing data in the long term.

## Unreliability of Location-based Identifiers

The problems related to using URLs for referencing datasets are largely due to the fact that they are location-based identifiers they describe where the data is but not necessarily what it is. Also, by definition, data accessed via URLs must be mediated by a central authority, such as the institutional repositories that serve biodiversity datasets, who can match location-based identifiers with data. Interested users are expected to trust the central authority to guarantee long-term access to the referenced data in its original form.

The use of URLs as identifiers violates the requirements of uniqueness and persistence (Paskin 1999). An identifier must only ever identify one entity (uniqueness) and must persist longer than the entity it identifies (persistence) (Paskin 1999). However, as we have shown in our experiments, many URLs do not possess both uniqueness and persistence; unstable URLs forfeit uniqueness in the event of content drift, while unresponsive URLs do not persist as long as the datasets they identify.



At the core of URL instability is the current practice of using URLs to 369  
identify evolving datasets rather than fixed dataset versions. If biodiversity 370  
data providers were uniformly committed to allocating one URL per 371  
dataset version, then content drift might become less common, improving 372  
overall URL stability; however, widespread social adoption of such a 373  
commitment from all data providers may be unrealistic. Additionally, such 374  
a commitment would not address link rot and URL unresponsiveness. 375  
Even if a similar commitment were made by data providers to guarantee 376  
the long-term responsiveness of URLs, it could not address the case where 377  
a data provider either loses authority over a domain name or migrates to 378  
another. For example, our deployed Preston observatories recorded the 379  
sudden migration of over 14,000 Plazi datasets from the 380  
`http://plazi.cs.umb.edu/` domain to `http://tb.plazi.org/`, an event which 381  
invalidated any references to URLs within the first domain. 382

Paskin proposed that “the best way to ‘future proof’ an identifier 383  
scheme is to forego any intelligence within the identifier itself” (Paskin 384  
1999), where the notion of intelligence refers to the inclusion of meaningful 385  
information in the textual representation of the identifier. URLs are 386  
typically structured according to the Domain Name System specification 387  
(though URLs may include an IP address instead of a domain name) and 388  
inherently contain some minimum amount of intelligence: the domain in that 389  
the URL belongs to (Mockapetris 1987). Thus, it is necessary to look to 390  
another identification scheme to allow for proper identification and reliable 391  
referencing. 392

## An Alternative: Unique Content-Based Identifiers

Instead of identifying digital datasets by location (e.g., a URL), we can identify datasets by their content. One way to achieve this is to use algorithmically generated content-based identifiers. A variety of cryptographic hashing algorithms are available which guarantee a unique hash, representable as text, for any given dataset (NIST 2001). Because the hash is deterministically derived from the content it identifies, we say that it is a content-based identifier. Because hashes are deterministic, anyone interested in identifying a dataset can simply compute its hash without the need for some mediating central authority (Paskin 1999). If a change is made to the dataset, then the hash computed from the modified dataset will be different from that of the original. Therefore, if the hash of a dataset is the same as the referenced hash, it must be the originally referenced dataset (figure 4(c)) (NIST 2001). Because hash identifiers can only identify the exact content that was referenced, content drift is impossible; a content hash will never match with either a different version of the content any other content.

The shift from location-based to content-based identifiers decouples future dataset accessibility from the original point of access. As long as there exists some discoverable and accessible data repository that serves the desired content can always be retrieved. Such data repositories can be made discoverable through content hash registries such as hash-archive.org (Trask 2015). In response to a user query for a content hash, these content hash registries would provide a list of locators (e.g., URLs), if any, that direct users to the referenced data (e.g., a registry would provide URLs that retrieve data when queried). Even if one

repository becomes inaccessible due to either a temporary outage or  
permanent retirement, another may be available to provide the referenced  
data. When several repositories serve referenced datasets, there is no single  
point of failure for content hash lookups; if a referenced dataset is  
redundantly located across and within data repositories, access to the  
dataset will only be lost if all associated locations exhibit link rot. Even if  
access to a dataset is lost, it can be restored as long as the referenced  
dataset still exists somewhere and can be made discoverable and accessible.

It is worth noting that the duplication of content across different  
information platforms does not lead to ambiguous references, but rather to  
distributed copies of the same reliably addressed content.

## Transitioning to Reliable References

Although we propose a change in the fundamental mechanisms used to  
reference datasets, existing references can be made reliable with only minor  
modifications. Consider the following citation generated by GBIF  
according to their citation guidelines (GBIF 2019):

Levatich T, Padilla F (2017). EOD - eBird Observation  
Dataset. Cornell Lab of Ornithology. Occurrence dataset  
<https://doi.org/10.15468/aomfnb> accessed via GBIF.org on  
2018-09-02.

The citation references the eBird dataset hosted at gbif.org as it was  
retrieved on September 9, 2018. However, at the time of writing, the URL  
<https://doi.org/10.15468/aomfnb> redirects to a GBIF internal reference  
page which states that the eBird dataset was last updated in March of

2019. The dataset made available through the listed URL is different from  
what was originally referenced in the citation, but it is impossible to  
determine the extent of the changes without having access to previous  
versions of the data.

Fortunately, references like the example above can be made more  
reliable by augmenting them with a content-based identifier for the dataset.  
Consider the following enriched citation for the eBird dataset that adds a  
SHA-256 content hash (NIST 2001):

Levatich T, Padilla F (2017). EOD - eBird Observation  
Dataset. Cornell Lab of Ornithology. Occurrence dataset hash:  
//sha256/29d30b566f924355a383b13cd48c3aa239d42cba0a55f4  
ccfc2930289b88b43c accessed at  
<https://doi.org/10.15468/aomfnb> via GBIF.org on 2018-09-02.

The content hash is captured in a content address Uniform Resource  
Identifier (URI) (Berners-Lee et al. 2005) in the form of  
hash://algo/hash-string proposed by (Trask 2015), where “algo” is a  
hashing algorithm (e.g., “sha256”) and “hash-string” is the content hash  
generated by the algorithm in hexadecimal format. In the example above,  
the hashing algorithm is SHA256 and the hash string starts with “29d  
The added content hash was derived from and uniquely identifies the exact  
version of the eBird dataset that was originally referenced. If an interested  
user knows of and has access to an information retrieval system that has  
indexed the dataset, finding the desired dataset is as simple as querying for  
its content hash. With the addition of a content hash, the URL becomes  
superfluous and is included merely to demonstrate that the URL and  
content hash are not mutually exclusive (see figure 4(b)).

Note that different hashing algorithms will generate different content hashes from the same data. We use a URI rather than the content hash itself because it allows us to specify the hashing algorithm used to generate the hash. If the hashing algorithm is not specified, one might mistakenly conclude that a dataset does not match a reference if the wrong hashing algorithm is used to verify the dataset’s authenticity. Our proposal to use Trask’s content-addressed URIs to reliably reference data is similar to, and was inspired by, Kuhn & Dumontier’s method to make digital content verifiable and permanent using *trask* URIs (Kuhn and Dumontier 2015). We chose to use Trask’s content hash URIs because they are location- and content-agnostic and easy to read. However, we recognize that *trask* URIs can help facilitate content retrieval and processing using a location-based URI prefix and an (optional) extension suffix *res*ively.

Content-based identification schemes exist that resist changes in format in digital content. For example, the universal numeric fingerprint (UNF) (Altman and King 2007) resists such changes by first preprocessing the input data before generating a content hash. Among other preprocessing techniques used when generating UNFs, numerical data may be rounded to a certain precision before generating a content hash, with the understanding that a dataset may undergo such format changes when translated, for example, between different computing environments or hardware configurations. Indeed, on manual examination of the changes between successive versions of the biodiversity datasets we observed, we found that, in the cases, two versions of a dataset (determined to be different because they resulted in different content hashes) differed only in formatting, such as the amount of whitespace and the sequential ordering

of observational records. However, for biodiversity data, we expect that such format-specific content-based identification schemes would only prove detrimental in practice. Standard cryptographic hashing algorithms such as SHA-256 are included in most modern software environments and enjoy widespread use across different digital applications, whereas non-standard algorithms such as UNF would first need to be installed and may be unknown to most users, presenting a hurdle to their widespread adoption. Additionally, it may be unrealistic to expect preprocessing efforts to filter out non-informative data effectively enough to be able to trust that semantically identical datasets will always result in the same content-based identifiers. This is especially relevant to biodiversity datasets because they consist mostly of text data (whereas UNFs were designed to hash primarily numerical data), which may be altered in a seemingly infinite number of ways without changing the actual content.




## Enhancing Dataset References with Provenance





A dataset reference can be given enhanced context by also referencing the record that describes its provenance. The following citation further augments the eBird dataset reference with the content hash of an associated provenance record:





Levatich T, Padilla F (2017). EOD - eBird Observation Dataset. Cornell Lab of Ornithology. Occurrence dataset hash: [//sha256/29d30b566f924355a383b13cd48c3aa239d42cba0a55f4ccfc2930289b88b43c](https://sha256/29d30b566f924355a383b13cd48c3aa239d42cba0a55f4ccfc2930289b88b43c) accessed at <https://doi.org/10.15468/aomfnb> via GBIF.org on 2018-09-02 with provenance hash: [//sha256/b83cf099449dae3f633af618b19d](https://sha256/b83cf099449dae3f633af618b19d)

05013953e7a1d7d97bc5ac01afd7bd9abe5d.

520

As was the case for the dataset, the provenance itself can be retrieved 521  
by querying an information system that has indexed the hash of the 522  
referenced provenance record. Note that the provenance hash is not 523  
strictly necessary to make a dataset reference reliable; the dataset hash 524  
alone is sufficient. However, explicitly referencing the provenance of the 525  
dataset is useful because it allows future readers to  retrieve the same 526  
context  that the original researcher  referenced the dataset had access 527  
to. More generally, the provenance describes the context of the retrieval of 528  
any type of content (e.g., datasets, metadata, citation files, etc.). The 529  
types of information in the provenance depend on the implementation of 530  
the data observatory, but at a minimum include the URLs that were 531  
queried to produce the content, the dates of the queries, the format of the 532  
content, and the data registries that were searched to find the content. 533

The use cases for  the  provenance hash are many. For example, 534  
if the provenance record of a dataset is found, it may be possible to 535  
traverse the provenance and find newer versions of the dataset. This 536  
requires that the various versions of the dataset were observed  at some 537  
 point in time by a provenance-generating data observatory, properly 538  
archived, then made publicly accessible. 539

A provenance record relates to a dataset the way that a map relates to 540  
a location: a provenance record provides a context to understand the 541  
origin and relations of a dataset. This provenance context may be limited 542  
to a  small amount  of meta-data elements related to a single dataset (e.g., 543  
web location, data format, author, license), but can also include a 544  
comprehensive description of a biodiversity dataset network  that  consists 545

of thousands of datasets and their associations. Also, because provenance records are datasets themselves, they can be reliably referenced and embedded in other provenance records using their content URIs. We used such a composition of content URIs and provenance records as part of our monitoring scheme (Jorrit Poelen and Patel 2018) to track the reliability of URLs in biodiversity data networks over time (see table 1 and figures 2 and 3). The following citation references the history of the entire DataONE network over the period of observation by one of our Preston observatories:

Poelen JH. 2019d. A biodiversity dataset graph: DataONE. doi:10.5281/zenodo.3483218.  
hash://sha256/87de0898919d2212977a586965e930ae45bdd1366073591c898c208a635e2814

## Dataset Retrieval Using Hash References

The dataset and provenance hashes referenced in the sample references above were produced by our Preston observatories which were set up to monitor the four data networks. Both the referenced dataset and its provenance are available online at zenodo.org (Poelen 2019c,d,f) and archive.org (Poelen 2019b). A query for the provenance hash in the search bar at zenodo.org or hash-archive.org should direct the user to an archived repository of Preston observations that contains both the dataset and its provenance (see figure 5). Given Zenodo’s long-term guarantee for data persistence and version availability (Zenodo 2019), the dataset reference is now reliable; it is effectively immune to both link rot and content drift. Future readers can trust that the dataset will stay available and, when downloaded, identically match the version of the eBird dataset we referenced. Note that, to comply with Zenodo’s limitations on user uploads




(Zenodo 2019), we only exposed the set of provenance hashes collected by  
each deployed Preston observatory for search indexing, which are far fewer  
in number than the dataset hashes. Thus, a query to zenodo.org for the  
dataset hash above should not produce any results. This is an artificial  
limitation; ideally, an information system would index the dataset hashes  
as well. Note that our Zenodo publication for the GBIF/iDigBio/BioCAsE  
observatory (Poelen 2019f) contains only provenance, although the Internet  
Archive publication (Poelen 2019b) contains the content as well as  
provenance. Our Zenodo and Internet Archive publications for BHL  
(Poelen 2019a,c) and DataONE (Poelen 2019d,e) contain both content and  
provenance.

Several biodiversity data aggregators, such as GBIF and iDigBio,  
produce a citation file for each user query to allow researchers to simply  
reference a single citation file rather than each individual dataset (GBIF  
2019, iDigBio 2016). A citation file lists the URLs of the datasets (among  
other things, such as attributions and retrieval dates) that were retrieved  
by the issued query. We have demonstrated that dataset URLs are  
unreliable references; thus, citation files that rely on URLs as references  
are also unreliable. Citation files could be made reliable if they were  
augmented with the hashes of the retrieved datasets and, optionally, their  
provenance records. In fact, citation files themselves can be referenced by  
hash, along with accompanying provenance hashes, as long as they are  
archived and made accessible.

## DOIs for Datasets and Queries

594

Biodiversity data aggregators often assign each dataset or query a Digital  
Object Identifier (DOI) (Paskin 2009) (e.g., 10.123/456) wrapped  URL  
(e.g., <https://doi.org/10.123/456>) and advise researchers to reference the  
generated DOI rather than a URL. Unfortunately, this abstraction does  
little to enhance the reliability of the reference.


595

596

597

598

599

The DOI Handle System (Paskin 2009) associates DOIs with online  
resources. However, it does not enforce any constraint  type of resource  
associated with a DOI. When DOIs are used to reference biodiversity  
datasets, the associated resources are often URLs, and therefore the use of  
such DOIs as referencing mechanisms is just as potentially unreliable as  
using URLs. In practice, these DOIs identify the evolving dataset (or set  
of datasets in the case of a query) rather than a fixed version, as  
demonstrated in the example references above. It is possible that an  
author would wish to make such a reference to an evolving online digital  
object. For example, an author promoting use of a published dataset might  
want future users to be directed to the most up-to-date content. However,  
such a fluid reference is not appropriate for making published results  
reproducible.

600

601

602

603

604

605

606

607

608

609

610

611

612

The Handle System allows for a complex web of redirection and  
distributed responsibilities. Just as the Domain Name System resolves  
domain names in URLs to IP addresses, the DOI Handle System allows  
DOIs to be resolved to URLs. However, the responsibility for resolving  
DOIs to URLs is divided between the Handle System and DOI registrars.  
The Handle System serves as the central authority that maps DOI prefixes  
to DOI registrars, examples of which include BHL, DataONE, GBIF, and

613

614

615

616

617

618

619

iDigBio. These registrars are then responsible, and index the central  
authorities for, associating DOIs that match their designated prefix with  
URLs, and are free to change the URL associated with any given DOI  
under their jurisdiction ( [IDF] International DOI Foundation 2018, Paskin  
2009).

The ability of biodiversity data networks to change the URL associated  
with a DOI is good for reference reliability in the sense that networks can  
account for dataset migration without compromising existing references.  
However, the use of DOIs addresses neither the instability of the URLs  
they redirect to nor cases of link rot in which no URLs remain responsive  
to serve the referenced dataset. Additionally, as the number of datasets  
identified online continues to grow, proper maintenance of all of the DOIs  
a data network administrates might become more unsustainable over time,  
potentially increasing the risk of unreliable URLs going undetected.

In an article proposing HTTP-URI-based stable identifiers (e.g., URLs  
that are resolvable over HTTP) for biological collection objects, Güntsch et  
al. admit that the use of DOIs does not solve the problem of unreliable  
referencing but merely deflects the burden of URL maintenance onto  
institutional repositories (Güntsch et al. 2017). In contrast, we propose a  
dataset referencing scheme that is reliable and can be supported by existing  
infrastructures and workflows. If existing workflows require references to  
be in the form of DOIs, it could be convenient to embed content hashes  
into DOIs. Such an approach has already been established for ISBNs  
through the creation of actionable ISBNs, or ISBN-As (Weissberg 2008),  
which may serve as a model for actionable content hashes.

## What It Means to Preserve Data

Our results indicate that reference rot poses an existential threat to published biodiversity datasets. We have seen that the use of content-based identifiers can effectively address the issue of reference rot. However, identifiers are of little use in a vacuum. An identifier can only be useful for data retrieval when combined with a resolver to associate identifiers with locations and a database to retrieve the dataset at the associated location (Paskin 1999). Thus, we need to address how resolvers and databases might be organized to accommodate content-based identifiers in order to fully realize long-term data preservation. In this context, we define data preservation as the continued capacity for datasets to be reliably referenced and retrieved in their original form even as the global digital biodiversity network evolves over time.


We propose four requirements that must be met to ensure proper data preservation that prevents data loss: 1) datasets must be addressable and retrievable using content-based rather than location-based identifiers; 2) an agent must exist to collect datasets, record their provenance, and deposit both to a dedicated repository; 3) these repositories should archive data that could be used in the future; and 4) content hash registries should be openly accessible to resolve hash identifiers to dataset locations within such repositories. Although openly accessible registries should make archived data discoverable, access to those data can still be restricted. Additionally, for the purposes of archival, it is important that the recorded provenance records do not necessarily describe the datasets themselves, but rather the activities that led to the procurement of those datasets; the primary purposes of provenance in the context of an archive are to

document the fact that evidence (i.e., an observation of a dataset) does  
exist and to make it discoverable for interested users (Bearman 1995).

We have shown that software agents such as Preston can be used to  
collect datasets and their provenance over time while maintaining  
content-addressability; all that is needed to ensure proper data preservation  
are a dedicated repository and an openly accessible content hash registry  
to map content-based identifiers to datasets located in the repository. In  
practice, repositories and registries (and potentially software agents such as  
Preston deployments) can be constructed; examples include Zenodo and the  
Internet Archive, although they impose some limitations that may restrict  
file size, number of files, and the amount of information that can be  
indexed (Internet Archive 2019, Zenodo 2019). These existing information  
systems may serve as models for long-term biodiversity information  
system. The requirements help to ensure that biodiversity data remain  
FAIR (Findable, Accessible, Interoperable, and Reusable) (Wilkinson et al.  
2016). Findability is achieved through the publishing of provenance logs  
which thoroughly describe what datasets are and where they originated  
from. The amenability of the content-based identification paradigm to the  
operation of independent decentralized repositories strengthens  
accessibility by preventing the failure of a single data repository from  
inhibiting future data access (see figure 4). Content-based identification  
also contributes to interoperability across data networks due to the  
absence of any central authority to administrate data access; a content  
hash computed from a dataset is guaranteed to match the hash computed  
by any other agent using the same dataset. Furthermore, content-based  
identifiers can be embedded in or referenced by DOIs to maintain

compatibility with systems that use DOIs as identifiers. Finally, and  
particularly relevant to this paper’s purpose, reusability is strengthened by  
enhancing the retrievability of referenced datasets and allowing users to  
verify that a retrieved dataset exactly matches that which was referenced.

## Conclusions

Although reference rot is resulting in a steady decline in the reliability of  
our digital biodiversity record, realistic solutions are available to address  
the root causes of the issue. Content drift can be eliminated altogether by  
changing the way we reference datasets from using location-based  
identifiers to ones that are content-based. Meanwhile, the online  
biodiversity data networks can be made more resilient to link rot if  
decentralized observation, val, and distribution techniques are used to  
capture incremental changes to the data record so that references can  
remain valid even when online datasets are updated, removed, or relocated.  
The use of content-based identifiers should be considered by biodiversity  
data aggregators in order to increase the reliability of references to the  
data they aggregate.

We have demonstrated that data observatories can be deployed to track  
the growing digital biodiversity data record. Using the dataset provenance  
collected over a period of seven months, we were able to quantify the  
change in reliability over time in terms of link rot and content drift  
exhibited by the URLs registered in major biodiversity data networks.  
Even if data networks uniformly adopted content-based identification of  
datasets and maintained versioned datasets, our method of quantifying link  
rot and content drift in data networks could be used to monitor whether

either of these issues persist in practice due to implementation flaws or  
non-technical issues.

Biodiversity data observatories can also be used by to increase the  
longevity of the biodiversity data record. Such observatories can be used  
to form reliable dataset references as well as recover datasets that would  
otherwise become inaccessible due to link rot and content drift.  
Additionally, the dataset provenance captured by such observatories serves  
as evidence of the evolution and distribution of the digital biodiversity  
data record. The combination of archived datasets and provenance can  
ensure the long-term reproducibility of scholarly works that reference  
ever-evolving biodiversity datasets.

Furthermore, the establishment of dedicated data repositories and  
publicly accessible content hash registries are beneficial for making  
content-addressed biodiversity data discoverable, distributable, and  
long-lived, by making the datasets and provenance captured by  
biodiversity data observatories securely archived and publicly available.

Great care has been taken to establish rigorous preservation guidelines  
for physical specimens, yet there is much that can be done to increase the  
longevity of our digital data. Our method is not only suited for tracking  
datasets in biodiversity data networks, but also provides a resilient and  
reliable way to publish, reference, and preserve scientific digital datasets  
without having to abandon our existing infrastructures. The method  
provides a much-needed foundation for constructing digital provenance  
graphs from an accessible, verifiable, and citable digital scholarly record.

## Acknowledgments

746

The research reported in this paper was funded in part by a grant (NSF 747  
OAC 1839201) from the National Science Foundation and the AT&T 748  
Foundation. The authors acknowledge early exchanges with Matt Collins, 749  
Anne Thessen, Jen Hammock, Katja Seltsmann, and Carl Boettiger. 750

## References

- [IDF] International DOI Foundation. 2018. Doi handbook. Technical report. International DOI Foundation. doi:10.1000/182. Accessed: 2019-12-04.
- Altman M, King G. 2007. A proposed standard for the scholarly citation of quantitative data. D-Lib Magazine 13.
- Bearman D. 1995. Archival strategies. The American Archivist 58:380–413. doi:10.17723/aarc.58.4.pq71240520j31798.
- Berners-Lee T, Fielding RT, Masinter LM. 2005. Uniform Resource Identifier (URI): Generic Syntax. Technical Report 3986. doi:10.17487/RFC3986. Accessed: 2019-12-04.
- Berners-Lee T, Masinter LM, McCahill MP. 1994. Uniform Resource Locators (URL). Technical Report 1738. doi:10.17487/RFC1738. Accessed: 2019-12-04.
- Costello MJ, Bouchet P, Boxshall G, Fauchald K, Gordon D, Hoeksema BW, Poore GCB, van Soest RWM, Stöhr S, Walter TC, Vanhoorne B, Decock W, Appeltans W. 2013. Global coordination and standardisation



in marine biodiversity through the world register of marine species (WoRMS) and related databases. PLoS ONE 8:e51629. doi:10.1371/journal.pone.0051629.

[DataONE] Data Observation Network for Earth. 2012. DataONE citation guidelines. <https://www.dataone.org/citing-dataone>. Accessed: 2019-12-04.

Davis EB, Schmidt D. 1996. Guide to Information Sources in the Botanical Sciences. Vol. 2nd ed. Reference Sources in Science and Technology. Englewood, Colo: Libraries Unlimited.

Edwards JL. 2000. Interoperability of biodiversity databases: Biodiversity information on every desktop. Science 289:2312–2314. doi:10.1126/science.289.5488.2312.

Elton CS. 1927. Animal ecology. Macmillan Co.,. doi:10.5962/bhl.title.7435.

Garfield E, Sher IH, Torpie RJ. 1964. The Use of Citation Data in Writing the History of Science. Institute for Scientific Information Inc Philadelphia PA.

[GBIF] Global Biodiversity Information Facility. 2019a. GBIF citation guidelines. <https://www.gbif.org/citation-guidelines>. Accessed: 2019-12-04.

[GBIF] Global Biodiversity Information Facility. 2019b. What is the GBIF? <https://www.gbif.org/what-is-gbif>. Accessed: 2019-12-04.

Güntsch A, Hyam R, Hagedorn G, Chagnoux S, Röpert D, Casino A, Droege G, Glöckler F, Gödderz K, Groom Q, Hoffmann J, Holleman A,

- Kempa M, Koivula H, Marhold K, Nicolson N, Smith VS, Triebel D. 2017. Actionable, long-term stable and semantic web compatible identifiers for access to biological collection objects. Database 2017. doi:10.1093/database/bax003.
- Hortal J, de Bello F, Diniz-Filho JAF, Lewinsohn TM, Lobo JM, Ladle RJ. 2015. Seven shortfalls that beset large-scale knowledge of biodiversity. Annual Review of Ecology, Evolution, and Systematics 46:523–549. doi:10.1146/annurev-ecolsys-112414-054400.
- [iDigBio] Integrated Digitized Biocollections. 2016. iDigBio citation guidelines. <https://www.idigbio.org/content/idigbio-terms-use-policy>. Accessed: 2019-12-04.
- Internet Archive. 2019. Uploading - a basic guide. <https://help.archive.org/hc/en-us/articles/360002360111-Uploading-A-Basic-Guide>. Accessed: 2019-12-04.
- Jorrit Poelen IA Michael Elliott, Patel P. 2018. Preston: a biodiversity dataset tracker. doi:10.5281/zenodo.1410543.
- Klein M, de Sompel HV, Sanderson R, Shankar H, Balakireva L, Zhou K, Tobin R. 2014. Scholarly context not found: One in five articles suffers from reference rot. PLoS ONE 9:e115253. doi:10.1371/journal.pone.0115253.
- Kuhn T, Dumontier M. 2015. Making digital artifacts on the web verifiable and reliable. IEEE Transactions on Knowledge and Data Engineering 27:2390–2400. doi:10.1109/tkde.2015.2419657.

- Matsunaga A, Figueiredo R, Thompson A, Traub G, Beaman R, Fortes JA. 2013. Integrated digitized biocollections (iDigBio) cyberinfrastructure status and futures.
- Michener W, Vieglais D, Vision T, Kunze J, Cruse P, Janée G. 2011. DataONE: Data observation network for earth: Preserving data and enabling innovation in the biological and environmental sciences. D-Lib Magazine 17. doi:10.1045/january2011-michener.
- Mockapetris P. 1987. Domain names - concepts and facilities. RFC 1034. doi:10.17487/RFC1034.
- [NIST] National Institute for Standards and Technology. 2001. Descriptions of sha-256, sha-384, and sha-512. <https://web.archive.org/web/20130526224224/http://csrc.nist.gov/groups/STM/cavp/documents/shs/sha256-384-512.pdf>. Accessed: 2019-12-04.
- Page LM, MacFadden BJ, Fortes JA, Soltis PS, Riccardi G. 2015. Digitization of biodiversity collections reveals biggest data on biodiversity. BioScience 65:841–842. doi:10.1093/biosci/biv104.
- Paskin N. 1999. Toward unique identifiers. Proceedings of the IEEE 87:1208–1227. doi:10.1109/5.771073.
- Paskin N. 2009. Digital object identifier (DOI®) system. In: Encyclopedia of Library and Information Sciences, Third Edition. CRC Press. p. 1586–1592. doi:10.1081/e-elis3-120044418.
- Pasquier T, Lau MK, Trisovic A, Boose ER, Couturier B, Crosas M,


- Ellison AM, Gibson V, Jones CR, Seltzer M. 2017. If these data could talk. *Scientific Data* 4. doi:10.1038/sdata.2017.114.
- Poelen JH. 2019a. A biodiversity dataset graph: Biodiversity Heritage Library (BHL). <https://archive.org/details/preston-bhl>. Accessed: 2019-12-04.
- Poelen JH. 2019b. Biodiversity Dataset Archive.  
hash://sha256/d79fb9207329a2813b60713cf0968fda10721d576dcb7a36038faf18027eebc1.  
<https://archive.org/details/biodiversity-dataset-archives>.  
Accessed: 2019-12-04.
- Poelen JH. 2019c. A biodiversity dataset graph: BHL.  
hash://sha256/67cc329e74fd669945f503917fbb942784915ab7810ddc41105a82ebe6af5482.  
doi:10.5281/zenodo.3484555.
- Poelen JH. 2019d. A biodiversity dataset graph: DataONE.  
hash://sha256/87de0898919d2212977a586965e930ae45bdd1366073591c808c208a635e2814.  
doi:10.5281/zenodo.3483218.
- Poelen JH. 2019e. A biodiversity dataset graph: DataONE.  
hash://sha256/87de0898919d2212977a586965e930ae45bdd1366073591c808c208a635e2814.  
<https://archive.org/details/preston-dataone>. Accessed:  
2019-12-04.
- Poelen JH. 2019f. A biodiversity dataset graph: GBIF, iDigBio, BioCAsE.  
hash://sha256/d79fb9207329a2813b60713cf0968fda10721d576dcb7a36038faf18027eebc1.  
doi:10.5281/zenodo.3484205.

- Poelen JH. 2019g. To connect is to preserve: on frugal data integration and preservation solutions. doi:10.17605/OSF.IO/A2V8G.
- Postel J. 1981. Internet Protocol. Technical report. doi:10.17487/RFC0791. Accessed: 2019-12-20.
- Rinaldo C, Norton C. 2009. BHL, the biodiversity heritage library: An expanding international collaboration. Nature Precedings doi:10.1038/npre.2009.3620.1.
- Robertson T, Döring M, Guralnick R, Bloom D, Wieczorek J, Braak K, Otegui J, Russell L, Desmet P. 2014. The GBIF integrated publishing toolkit: Facilitating the efficient publishing of biodiversity data on the internet. PLoS ONE 9:e102623. doi:10.1371/journal.pone.0102623.
- Trask B. 2015. Principles of content addressing. <https://bentrask.com/?q=hash://sha256/98493caa8b37eaa26343bbf73f232597a3ccda20498563327a4c3713821df892>. Accessed: 2019-12-04.
- Vision TJ. 2010. Open data and the social contract of scientific publishing. BioScience 60:330–331. doi:10.1525/bio.2010.60.5.2.
- Weissberg A. 2008. The identification of digital book content. Publishing Research Quarterly 24:255–260. doi:10.1007/s12109-008-9093-8.
- Wieczorek J, Bloom D, Guralnick R, Blum S, Döring M, Giovanni R, Robertson T, Vieglais D. 2012. Darwin core: An evolving community-developed biodiversity data standard. PLoS ONE 7:e29715. doi:10.1371/journal.pone.0029715.


Wilkinson MD, Dumontier M, Aalbersberg IJ, Appleton G, Axton M, Baak A, Blomberg N, Boiten JW, da Silva Santos LB, Bourne PE, Bouwman J, Brookes AJ, Clark T, Crosas M, Dillo I, Dumon O, Edmunds S, Evelo CT, Finkers R, Gonzalez-Beltran A, Gray AJ, Groth P, Goble C, Grethe JS, Heringa J, 't Hoen PA, Hooft R, Kuhn T, Kok R, Kok J, Lusher SJ, Martone ME, Mons A, Packer AL, Persson B, Rocca-Serra P, Roos M, van Schaik R, Sansone SA, Schultes E, Sengstag T, Slater T, Strawn G, Swertz MA, Thompson M, van der Lei J, van Mulligen E, Velterop J, Waagmeester A, Wittenburg P, Wolstencroft K, Zhao J, Mons B. 2016. The FAIR guiding principles for scientific data management and stewardship. *Scientific Data* 3. doi:10.1038/sdata.2016.18.

Zenodo. 2019. General policies. <https://about.zenodo.org/policies/>. Accessed: 2019-12-04.

## Figures




(a) Physical specimen collection




(b) Digital data collection


**Figure 1.** Reliable record keeping for digital datasets (b) can be achieved in an analogous way to current practices in record keeping for physical specimens (a). Biologists collect physical specimens from the natural world, thoroughly document the process, then store the specimens in facilities equipped for long-term preservation. Analogously, digital datasets that are downloaded from the internet can be thoroughly documented and archived in dedicated repositories for long-term preservation. Just as the collection of physical specimens is recorded and identified in specimen history records, the downloading of digital datasets can also be recorded and identified in dataset history records.




(a) BHL




(b) DataONE



(c) GBIF




(d) iDigBio




**Figure 2.** Overall responsiveness, stability, and reliability from March 2019 through October 2019 as percentages of URLs that exhibit each indicator in (a) BHL, (b) DataONE, (c) GBIF, and (d) iDigBio.







(a) BHL



(b) DataONE



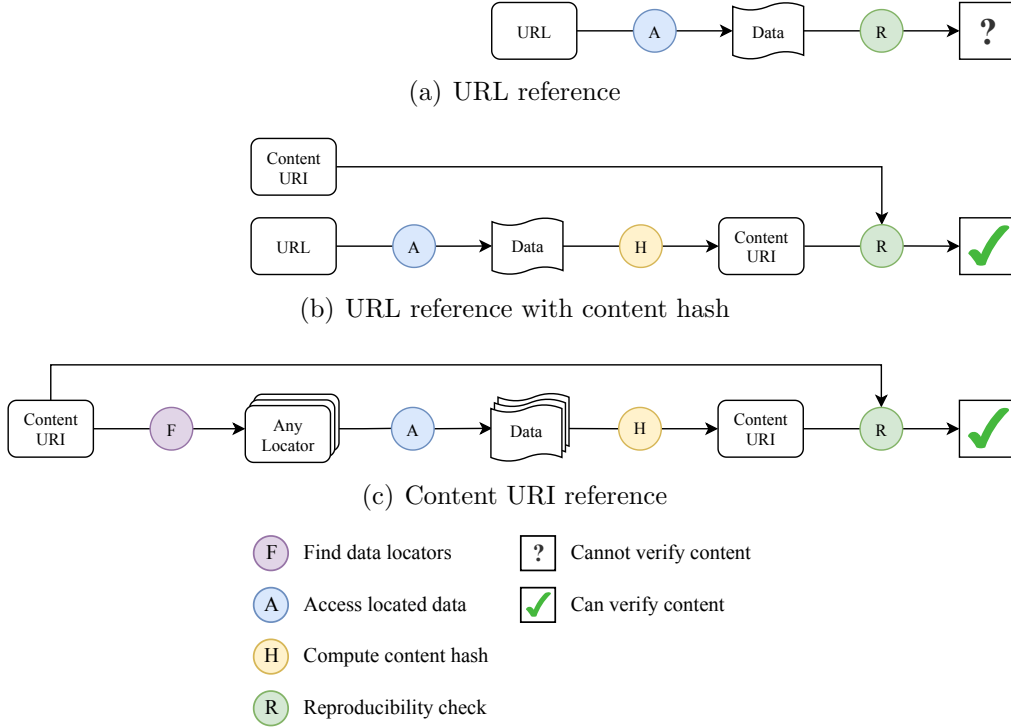
(c) GBIF



(d) iDigBio



**Figure 3.** Total number of URLs and unique contents observed from March 2019 through October 2019 in (a) BHL, (b) DataONE, (c) GBIF, and (d) iDigBio.



**Figure 4.** Content resolution and verification for references that use location- versus content-based identifiers. (a) Location-based identifiers (e.g. URLs) cannot verify the authenticity of retrieved content and are vulnerable to link rot due to the use of a fixed locator. (b) If the content hash of the referenced data is known, the authenticity of retrieved data can be verified by comparing the hash of the retrieved data with the provided content hash. However, the fixed locator is still vulnerable to link rot. (c) Content-based identifiers (e.g. Content URIs) can be used to find several locators for the referenced data and contain a content hash to verify the authenticity of retrieved data. The decoupling of the reference from a fixed locator makes the reference resistant to link rot.

**Hash Archive** (beta)

URL or hash:

Sources for [hash://sha256/b83cf099449dae3f633af618b19d05013953e7a1d7d97bc5ac01afd7bd9abe5d](https://archive.org/download/biodiversity-dataset-archives/data.zip/data/b8/3c/b83cf099449dae3f633af618b19d05013953e7a1d7d97bc5ac01afd7bd9abe5d/)

- [Search for this hash on Google](#)
- [Search for this hash on DuckDuckGo](#)
- [Search for this block on IPFS](#)
- [Check this hash on VirusTotal](#)
- [Other useful sources...?](#)

*Active as of November 5<sup>th</sup>, 2019*

<https://archive.org/download/biodiversity-dataset-archives/data.zip/data/b8/3c/b83cf099449dae3f633af618b19d05013953e7a1d7d97bc5ac01afd7bd9abe5d/><sup>[1]</sup>

*Active as of October 8<sup>th</sup>, 2019*

<https://deeplinker.bio/b83cf099449dae3f633af618b19d05013953e7a1d7d97bc5ac01afd7bd9abe5d/><sup>[1]</sup>

**Figure 5.** An example of a search index mapping hashes to archives. A search for a content or provenance hash at hash-archive.org will find any associated URLs that have been registered at hash-archive.org.

## Tables

<b>Data Network</b>	<b>Responsive URLs</b>	<b>Stable URLs*</b>	<b>Reliable URLs</b>
BHL <sup>a</sup>	57.41% (142,672)	99.97% (232,996)	57.39% (142,633)
DataONE <sup>b</sup>	94.55% (352,438)	92.27% (339,109)	87.09% (324,641)
GBIF <sup>c</sup>	71.72% (49,707)	37.35% (20,094)	24.05% (16,669)
iDigBio <sup>c</sup>	88.04% (5,477)	68.69% (4,251)	61.68% (3,837)
All observed URLs**	78.94% (546,645)	90.43% (593,469)	70.07% (485,203)

**Table 1.** Overall responsiveness, stability, and reliability for URLs observed in each biodiversity data network and for all observed URLs as of October 2019. \*URLs that never provided content were omitted from the divisor when calculating Stable URLs percentages. \*\*Because URLs may be registered in more than one data network, the total number of observed URLs is expected to be less than the sum of the URL counts for each network. <sup>a</sup>Poelen (2019c) <sup>b</sup>Poelen (2019d) <sup>c</sup>Poelen (2019f)